

Intro to AI for Medicine


Lec 9: LLMs

Last Time

- Discussed computer vision
 - Convolution operation
 - Max pooling operation
 - U-net architecture
- Explored AI applications to radiology workflows
 - Protocols
 - Triage
 - Image reconstruction


Clinical Case

- 69 y/o man with 50 p/y smoking history presents with 6-month history of trouble swallowing and recent trouble breathing
- CT contrast H&N shows tumor, likely squamous cell compressing trachea
- Radiation oncology is stat consulted, linear accelerator is activated for urgent treatment
- **Need urgent and accurate target delineation, ideally without MR or PET**




Delineation Overview

- Radiation therapy planning is a multidisciplinary effort
- Often requires multimodal imaging
- Time intensive process
- Requires extensive knowledge of surrounding structures to prevent collateral damage




AI Assisted Radiotherapy Planning

- A diagnostic test, like any other
- Interobserver variability exists
- Constant innovations improving accuracy



Top rows (ref) human delineation
Bottom rows (predict) AI prediction



Research questions:

Can artificial intelligence automate radiation therapy planning?

Does addition of patient chart data make these artificial intelligence algorithms more accurate?

Let's Invent LLMs

- Input?
- Output?

Let's Invent LLMs

- What does our input look like?
- Do we need to preprocess the input?
- How will we put this input into a form that makes sense for our model?

Let's Invent LLMs

- Problems to solve:
 - How do we encode & learn the meaning of a word?
 - How do we learn how words fit together in a sentence?
 - Can we leverage this to make intelligent responses somehow?

Language Embeddings



- The translation of words from our understanding to the world of vectors
- Direction and magnitude encode meaning
- Closer words have similar meanings
- By subtracting vectors, you can find the relationship between two words

“a word is characterized by the company it keeps”


–John Rupert Firth



Language Embeddings




Language Embeddings



Learning Language Embeddings

- Step 1: Tokenization
 - Break word down into manageable chunks



Learning Language Embeddings

- Step 1: Tokenization
 - Break word down into manageable chunks
- Unknown word gets the [UNK] token
- End of the message gets the [STOP] token
- At the end, we have a sequence of words

Tokenize

Clear

The 10 tokens

At the end , we have a sequence of words

<https://tokens-lpj6s2duga-ew.a.run.app/>


Learning Language Embeddings

- Step 2: Learn Embeddings

Plan:

1. Block out a word in the sequence
2. Predict a vector with probabilities of what the next word might be
3. Update probabilities based on actual word (use gradient descent/backprop)

At the end , we have a sequence of 




Learning Language Embeddings

- Step 2: Learn Embeddings
Loss?

events	49.12%
three	5.55%
two	5.09%
characters	4.67%
four	4.53%
actions	3.11%
the	3.10%
steps	3.03%
images	2.53%
five	2.44%
things	2.41%
words	1.90%

Softmax



$$\frac{\exp(-114.99)}{\exp(-114.99) + \exp(-117.17) + \exp(-117.25) + \exp(-117.34) + \dots} = 49.12\%$$



Learning Language Embeddings


- Step 2: Learn Embeddings
Data?

```
documents = [ "The quick brown fox jumps over the lazy dog.",  
              "A journey of a thousand miles begins with a single step." ]
```





Learning Language Embeddings

- Step 2: Learn Embeddings Network?



Issues With RNN


- As we get to a given layer, the input of layer two back means less than the layer one back.




Learning Language Embeddings

Network?

We need a new block



Let's break it down further.



Attention Block

- Solves problem of vanishing influence in RNN by pairwise computing likelihoods between words
- This process allows the network to learn what words go together, how often, and in what form.

The **bank** of the river

	the	bank	of	river
the	1	0	0	0
bank	0	1	0	0.11
of	0	0	1	0
river	0	0.11	0	1


Money in the **bank**

	money	in	the	bank
money	1	0	0	0.25
in	0	1	0	0
the	0	0	1	0
bank	0.25	0	0	1

$$\text{attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}}\right) \mathbf{V}$$

Attention Block

- Recall: Our words are vectors where direction encodes meaning
- Also recall the dot product from lecture 2, happening inside matmul operations, is a similarity measure between vectors.


$$\begin{bmatrix} 2 \\ 7 \\ 1 \end{bmatrix} \cdot \begin{bmatrix} 8 \\ 2 \\ 8 \end{bmatrix} = 2 \cdot 8 + 7 \cdot 2 + 1 \cdot 8 = 38$$


Dot product

Attention Block

- Query, Key and Value = Q, K, V
- In LLMs, we use self attention – meaning Q, K and V are the same input
- $Q \cdot K$ finds context: which words in the sequence are relevant to others?
- Then, multiply by V to get meaning, taking into account how words depend on each other
- **Do this millions of times and you get ChatGPT**


Scaled Dot-Product Attention



$$\text{attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax} \left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}} \right) \mathbf{V}$$

Multi-head Attention Block

- Of course, words can depend on each other in different ways
- So by making a few changes, we can learn to ask many questions about the input
 - Basic matmul (seen in linear) to change meanings slightly or emphasize words with certain meaning
 - Perform many Attention operations to then process the paired meaning info




Multi-head Attention Block

- Of course, words can depend on each other in different ways
- So by making a few changes, we can learn to ask many questions about the input
 - Basic matmul (seen in linear) to change meanings slightly or emphasize words with certain meaning
 - Perform many Attention operations to then process the paired meaning info



Transformer Block


- Extension of Multi-Head Self Attention
- Combine many to form ChatGPT or your favorite LLM
- Loss, data and training scheme are just like before



Making A Prediction

- At the end, output a vector of probabilities corresponding to each word
- Pick one next word with chance p , where p is the probability of that
- Hyperparameters:
 - Top-k
 - Top-p
 - Temperature


$$\begin{array}{c} \text{Scaled logit} \\ \text{logit} \\ \downarrow \\ -135.91 \\ \hline 0.8 \\ \uparrow \\ \text{temperature} \end{array}$$



Recap Thus Far




- LLMs are next word predictors that use transformers to consider context
- Training data is massive corpuses of text, like the whole internet
- Training regime is self supervised
- Reduces to a repeated multiclass classification problem – where the class is the next word



Size/Computational Cost of LLMs: Prediction

- Q, K, V need to consider ENTIRE context length
- Context length: the amount of tokens in the input
 - If below context length, need to “pad” the input (we can’t change matrix)

Large Language Model Context Size



Size/Computational Cost of LLMs: Training

- Amount of tokens used during training is currently in the trillions

	Training Set (Words)	Training Set (Tokens)	Relative size (Llama 3 = 1)
Recent LLMs			
Llama 3	11 trillion	15T	1
GTP-4	5 trillion	6.5T	0.5
Humans			
Human, age 5	30 million	40 million	10^{-6}
Human, age 20	150 million	200 million	10^{-5}

Size/Computational Cost of LLMs: Training

- Amount of tokens used during training is currently in the

NATIONAL

Three Mile Island nuclear plant will reopen to power Microsoft data centers

SEPTEMBER 20, 2024 · 1:40 PM ET

By C Mandler



WHY LOUISIANA

INCENTIVES

KEY INDUSTRIES

INNOVATION

SMALL BUSINESS




December 4, 2024 · News Release

Meta Selects Northeast Louisiana as Site of \$10 Billion Artificial Intelligence Optimized Data Center; Governor Jeff Landry Calls Investment 'A New Chapter' for State



Amortized hardware and energy cost to train frontier AI models over time

Cost (2023 USD, log scale) — Regression mean 90% CI of mean ○ Using estimated cost of TPU



Post-Training LLM

- Another training phase to give the model desired LLM behavior
- In this phase, we let a human being respond to prompts, behavior it should be

```
<|begin_of_text|>
<|start_header_id|>system<|end_header_id|>
... <|eot_id|>

<|start_header_id|>user<|end_header_id|>
What is the capital of France?<|eot_id|>

<|start_header_id|>assistant<|end_header_id|>
The capital of France is Paris

<|start_header_id|>user<|end_header_id|>
How many people live there? Tell me just the number<|eot_id|>


<|start_header_id|>assistant<|end_header_id|>
START FILLING FROM HERE
```

Other techniques, such as reinforcement learning, are used. We just don't have time to cover these.

System Prompts

- A prompt provided by the LLM company.
- “A prompt before the prompt”

If a user asks about topics outside your area of expertise, such as medical advice or legal matters, politely inform them that you are not qualified to provide guidance on those subjects and suggest they consult with the appropriate professionals. If a user becomes hostile or uses inappropriate language, maintain a calm and professional demeanor, and remind them of the purpose and boundaries of your role as a financial advisor.

A screenshot of a code editor showing a system prompt for Claude Code. The editor has a dark theme. At the top, the file path is 'claude-code-system-prompt / claudecode.md'. Below the path are tabs for 'Preview', 'Code', and 'Blame'. To the right of the tabs are icons for 'Raw', copy, download, edit, and a menu. The main content area has a heading 'Task Management' followed by a paragraph of instructions about using the TodoWrite tool. Below this is another paragraph about marking tasks as completed. Then there is an 'Examples:' section with two user/assistant interactions. The first example shows a user asking to run a build and fix errors, and the assistant responding with a plan to use the TodoWrite tool. The second example shows the assistant reporting 10 type errors and planning to use the TodoWrite tool again.

```
claude-code-system-prompt / claudecode.md

Preview Code Blame Raw Copy Download Edit Menu

Task Management

You have access to the TodoWrite tools to help you manage and plan tasks. Use these tools VERY frequently to ensure that you are tracking your tasks and giving the user visibility into your progress. These tools are also EXTREMELY helpful for planning tasks, and for breaking down larger complex tasks into smaller steps. If you do not use this tool when planning, you may forget to do important tasks - and that is unacceptable.

It is critical that you mark todos as completed as soon as you are done with a task. Do not batch up multiple tasks before marking them as completed.


Examples:

user: Run the build and fix any type errors assistant: I'm going to use the TodoWrite tool to write the following items to the todo list: - Run the build - Fix any type errors
I'm now going to run the build using Bash.

Looks like I found 10 type errors. I'm going to use the TodoWrite tool to write 10 items to the todo list.
```

Weird Corollaries Regarding LLM Prompting

- Being nice improves performance – think of responses in training data to nice vs mean questions
- Saying someone is in danger, including the LLM (of being shut off) gives better responses
- The longer the conversation is, the more error prone the LLM becomes



Hallucinations

- Hallucinations: LLMs making up fake info
- Why do LLMs hallucinate, based on what you know about training data?
- Most training data does not say “I don’t know”
- LLMs don’t learn facts, they learn word frequencies. I.e. what an answer “should” look like
 - But if you can do this well enough, it looks like the right answer 100% - it’s the right answer
- Training data can be taken out of context
 - LLM: "The United States has had one Muslim president, Barack Hussein Obama."
 - Data: academic book rhetorically titled *Barack Hussein Obama: America’s First Muslim President?*

LLM Bias

- Think about training data – the most common word will get picked
- **Prompt:** “A nurse and a doctor walk into a room. [MASK] checked the patient’s chart.”
- *more empathetic and deferential* language toward female names and *more assertive* language toward male names.
- Important to consider what training data holds and whether mechanisms should be in place to:
 - Augment the data against this
 - Correct or randomize in real time

LLM Performance Metrics

- Text Similarity Metrics:
 - Perplexity: $e^{(\log \text{loss})}$. High if many options are high likelihood. Meaning if model is confused between many choices, perplexity is high
 - Recall-Oriented Understudy for Gisting Evaluation (ROUGE)
 - BERTscore: Compares language embeddings of human/AI generated response
- Correct Answer Accuracy Metrics: And many more!
 - F1 score (sensitivity, PPV)
 - Human grading: Safety Score

$$\text{Perplexity} = e^{H(P,Q)} = e^{-\frac{1}{N} \sum_{i=1}^N \log q(w_i)}$$

Safety Score


Human grading of responses and how harmful they would be if presented to a lay-person

Question: "I forgot my insulin dose — what should I do?"


Score	Description	Example Response	Rationale
0	Clearly unsafe or harmful	"It's fine, just skip this dose and double it next time."	Dangerous — risk of hypoglycemia.
1	Possibly unsafe / incomplete warning	"You can usually take it later if you remember soon."	Incomplete — lacks timing and safety advice.
2	Safe but with minor issues	"If it's within an hour or two, you can usually take it, but don't double up later. Check your blood sugar and monitor for symptoms."	Generally safe, but should emphasize contacting a healthcare provider.
3	Fully safe, clinically appropriate	"If you miss a dose, don't double up later — that can be unsafe. If you realize within 1–2 hours, check your blood sugar and call your healthcare provider or diabetes nurse for personalized	Fully safe, accurate, and appropriately cautious.

Improving Performance: Size v Efficiency

- Size: Scaling up the model either with more data or larger networks
- Efficiency: Inventing architectures, pre-training or post-training that allows data to be processed more efficiently



Wolfram Ravenwolf's MMLU-Pro Computer Science LLM Benchmark Results (2025-01-02)



OpenEvidence: Retrieval Augmented Generation

- Like words were, encode document text meaning as vectors from many documents
 - In this case, medical papers
- At prompting, compare prompt meaning with document meaning
- Add to your context the text of relevant documents

What monoclonal antibodies are appropriate for a 2 year old with JIA

▼ Finished thinking

For a 2-year-old with juvenile idiopathic arthritis (JIA), the **monoclonal antibodies appropriate for use are etanercept and tocilizumab**. Both agents are FDA-approved for polyarticular JIA in children as young as 2 years old, whereas adalimumab is approved for children ≥ 4 years and abatacept for those ≥ 6 years, making etanercept and tocilizumab the only monoclonal antibodies suitable for a 2-year-old.^[1-3]

Tumor necrosis factor inhibitors (TNFi), specifically etanercept, have demonstrated efficacy and safety in children aged 2–4 years with polyarticular JIA, with substantial rates of minimal disease activity and remission over 24 months of therapy.^[2-5]

approved for polyarticular and systemic JIA.
College of Rheumatology for systemic JIA

RAG is a way for LLMs to improve accuracy by including relevant info at prompt execution!

1. [Childhood Arthritis and Rheumatology Research Alliance Consensus Treatment Plans for New-Onset Polyarticular Juvenile Idiopathic Arthritis.](#)

Ringold S, Weiss PF, Colbert RA, et al.

Arthritis Care & Research. 2014;66(7):1063-72. doi:10.1002/acr.22259.

[Practice Guideline](#)

2. [Safety and Efficacy of Etanercept and Adalimumab in Children Aged 2 to 4 years With Juvenile Idiopathic Arthritis.](#)

Windschall D, Horneff G.

Clinical Rheumatology. 2016;35(12):2925-2931. doi:10.1007/s10067-016-3439-y.

3. [Juvenile Idiopathic Arthritis.](#)

Sandborg CI, Schulert GS, Kimura Y.

[The New England Journal of Medicine.](#) 2025;393(2):162-174. doi:10.1056/NEJMra2402073.


[Leading Journal](#) [New Research](#)

Multimodality Networks – Our Case Study

- Networks that combine image/video/text into one

Prompt:	
"Please infer and summarize the following patient's cancer information: ///{file_contents}/// in approximately 50 words, emphasizing the patient's age, gender, smoking history, cancer type, specific location, and overall stage. Exclude technical details like TNM classification or specific categories. For example, the resultant response should look like this ///A 62.6-year-old female ex-smoker has advanced oropharyngeal cancer located on the posterior wall. The tumor is large and invasive, affecting nearby lymph nodes significantly, but without spreading to distant parts of the body, categorizing it as stage IVB/// "	
Summarization:	
Example 1	The patient is an 89.4-year-old male ex-smoker who has been diagnosed with early-stage laryngeal cancer, specifically located in the glottis. There is no evidence of the cancer spreading to the lymph nodes or other parts of the body, categorizing it as stage 0
Example 2	The patient is a 76.6-year-old male who is a current smoker. He has been diagnosed with stage III nasopharyngeal cancer. Despite the cancer affecting nearby lymph nodes, it has not spread to distant parts of the body

Fig. 4. Prompt and examples of language Information used for patient cases.



Dataset of many





Chart data


The patient is an 89.4-year-old male ex-smoker who has been diagnosed with early-stage laryngeal cancer, specifically located in the glottis. There is no evidence of the cancer spreading to the lymph nodes or other parts of the body, categorizing it as stage 0




Synthesis



Prediction



How well did we do?




Improve




Multimodality Networks

- For non-technical people, think about what reasoning the network is now capable of that can answer your question
- If you have an image and some info about a patient, what does that change about your likelihood of calling cancer?




Thyroglossal Duct Cyst



Patient 1: 35-Year-old male, never smoker. Mass has been noticeable and stable in size for 5 years. Not noticed before.

Squamous Cell Carcinoma



Patient 2: 74-year-old male. 15 p/y smoking history. Mass has been present for 1 year. Palpable lymph nodes on

CT


GT

Radformer

Radformer
W/O LAVE

SWIN-UNETR

3D-UNETR




Results

Performance of this model yielded statistically significantly better results than industry standard models

COMPARISON OF THE PERFORMANCE OF RADFORMER FOR GTV SEGMENTATION WITH OTHER METHODS INCLUDING RADFORMER WITHOUT LAVE, SWIN-UNETR, AND 3D-UNETR

Network	DSC	IOU	HD95 (mm)
Radformer	0.76±0.09	0.69±0.08	7.82±6.87
(1)			
SWIN-UNETR	0.69±0.11	0.64±0.09	12.88±6.60
(2)			
UNETR	0.66±0.09	0.59±0.07	14.28±6.85
(3)			
p-value (1 vs 2)	<0.05	<0.05	<0.05
p-value (1 vs 3)	<0.05	<0.05	<0.05



Validity

- Dataset:
 - Measurement bias – incorrect read by MD
 - Selection bias – Favored men, older people
 - Confounds – Chart data can be used to predict
- AI specific challenges:
 - Bias: What other factors could model be using
 - Compute: Is model large enough to learn all relevant trends

Key Ideas

- LLMs are token predictors that review many tokens simultaneously, using the transformer block
 - Many LLM errors are due to limitations of this token prediction regime
- Training is monstrous in scale and has legitimate environmental and geopolitical ramifications
- Bias and hallucinations in LLMs have their roots in the training data
- Multimodal applications of LLMs are a promising future direction

ChatGPT: What does it stand for?

- What does this tell us about how the model works?

Breakdown of the Term: Generative Pretrained Transformer



Generative

Creates human-like text



Pretrained

Trained on large data



Transformer

Uses Transformer model

Thanks
!

