

Intro to AI for Medicine

Lec 14: Ethical Issues in AI

Last Time

- Discussed sources of bias and how they can be worsened by AI
- Went over some of the classes of bias
- Had case studies and discussions regarding two instances where AI was biased

This time


- Moral and ethical dilemmas when it comes to AI
- Last time: how were models making biased decisions
- This time: is the decision making process fair to begin with
- Too many areas to touch on, so we'll only cover a handful

This Time

- Exploring the ways AI algorithms exhibit bias, harmful and otherwise
- Let what you know about AI algorithms inform your opinions
- Lots of discussion! You guys know more than me
- Topics can be heavy so let's tackle them with an open mind




AI Workflow



Asimov's Three Laws of Robotics

- The Three Laws, presented to be from the fictional "Handbook of Robotics, 56th Edition, 2058 A.D.", are:[\[1\]](#)
 - A robot may not injure a human being or, through inaction, allow a human being to come to harm.
 - A robot must obey the orders given it by human beings except where such orders would conflict with the First Law.
 - A robot must protect its own existence as long as such protection does not conflict with the First or Second Law.



Ethical Issues Regarding Data: Data Ownership

- Data collected by health systems, stored in their EHR
- From there, the EHR company or healthcare system is free to make AI products using this data
- Do the patients whose data is being used deserve some degree of compensation?

The importance of these drivers is illustrated by the variation in commercial value observed for various patient records.¹ For example, episodic patient data can sell for approximately \$75 per record, while an individual's genomics data can sell for about \$1,300 per record. Each of these patient records may be worth **two to three times more**, depending on levels of quality, robustness, detail or degree of integration with other datasets.

Database Credentialed Access

MIMIC-III Clinical Database

Alistair Johnson , Tom Pollard , Roger Mark 

	Patient care	Commercial operations	Regulatory compliance
Life sciences and medical device companies	• Research and development - Biomarker/target discovery and validation - Protocol design, synthetic control arms - Patient stratification and recruiting - CDx development • Medical affairs - Medical education - Comparative effectiveness	• Value analysis - Pricing strategy (optimization and justification) - HEOR - Payer engagement and reimbursement support • Commercial marketing - Brand analytics and commercial effectiveness - Competitive intelligence - Patient engagement	• Market access • Pharmacovigilance (e.g., post-market surveillance)
Payer	• Comparative effectiveness • Benefit/product design support	• Patient engagement (e.g., patient education, decision support) • Care management benchmarking/oversight • Price transparency	• Fraud, waste and abuse analysis
Providers	• Treatment patterns • Comparative effectiveness • Clinical decision support	• Competitive benchmarking • Internal process improvements	N/A
Regulators	• Data management, remediation and enhancement • Disease landscape/epidemiology • SDOH-based risk assessments • Program impact analytics (e.g., health/clinical outcomes)	N/A	• Fraud, waste and abuse analysis • Beneficiary identity assurance and eligibility • Provider compliance/credentialing • Pharmacovigilance
Healthcare data	• Interoperability/data integration • Clinical trial optimization	• Price transparency • Patient engagement	• Compliance support

Decentralized data ownership

- Instead of one organization owning data, smaller entities own and lease data to other parties
- More equitable in terms of ownership and compensation
- However, decreases standardization and adds red tape to datasets being assembled
- What do you all think about this model, or how could it be changed to better represent all stakeholders?

Ethical Issues Regarding Data: Privacy

- The obvious: data must be kept private and safe to respect patient confidentiality
- When too far out of context, ChatGPT shown to regurgitate training data verbatim
- Could happen with medical data as well





Figure 1: The map from prompts to latent space is injective. SIPIT inverts it.

Inference: Likeness Ownership

- Can make AI generated pictures of well known characters like Spongebob – but what if it's a real person?
- Who should own a likeness human or otherwise?
- Is placing data on the internet consenting to have it used in AI models?
- What other issues are created by generating images of others without their permission?




Medical Corollary – Digital Twins



- A simulated version of a patient that allows different treatments or interventions to be tested ahead of time
- Generated using a combo of AI and regular algorithms
- Ex. A model of heart rhythm which captures how currents travel in the myocardium
- Do patients own their digital twin for further experimentation, even if deidentified?
- What other issues might outweigh the benefits of digital twins?

AI Alignment Problem

- AI's generated responses should match the morals and ethical principals of its creators
- Entire divisions of large AI companies devoted to alignment
- Massive consideration wrt the problem of superintelligence – if AI is more powerful than us it should have our best interests in mind



Reward Hacking

AI Alignment Problem

- Forward alignment: alignment through training
 - Data selection
 - Post-training (in the case of LLM)
 - Preprompting
- Backward alignment: identifying and diagnosing issues in final model before or after deployment


Example 1: Instruction-Following Issues


Prompt: "Write a short poem about the ocean, focusing on its calmness, without mentioning any colors."

Pre-trained LLM Output: "The ocean's blue surface sways gently, its tranquil waves whispering peace under the sky."

AI Alignment Problem: Human in the loop

Reinforcement Learning from Human Feedback (RLHF)





<https://arxiv.org/pdf/2310.19852>

Figure 10: A Tree diagram summarizing the key concepts, logic, and literature related to Safety Evaluation. The root of the tree represents Safety Evaluation, which aims to *measure the accidents caused by design flaws in AI systems and harmful events that deviate from the intended design purpose of the AI system*. The main branches represent the main structure of safety evaluation, including Datasets and Benchmarks, Evaluation Targets, and Red Teaming techniques. Further sub-branches list key works exploring each of these branches. This diagram provides an overview of research directions and specific techniques for measuring AI systems' safety alignment degree.

Inference: Hallucinations

- Hallucinations: When AI presents information as fact that isn't true
- Recall
 - Training data seldom says "I don't know"
 - The objective of AI is to predict the next most accurate word, not create factually correct info

Here's What Happens When Your Lawyer Uses ChatGPT

A lawyer representing a man who sued an airline relied on artificial intelligence to help prepare a court filing. It did not go well.

 Share full article



 1.1K

Volume 19, Issue 6, June 2024, Pages 2106-2111



Successful management of an Iatrogenic portal vein and hepatic artery injury in a 4-month-old female patient: A case report and literature review

Show more ▾

+ Add to Mendeley Share Cite

<https://doi.org/10.1016/j.radcr.2024.02.037>

[Get rights and content](#) ↗

Under a Creative Commons [license](#) 

● open access

ChatGPT?



Surfaces and Interfaces

journal homepage: www.sciencedirect.com/journal/surfaces-and-interfaces

The three-dimensional porous mesh structure of Cu-based metal-organic-framework - aramid cellulose separator enhances the electrochemical performance of lithium metal anode batteries

Manshu Zhang^{a,1}, Liming Wu^{a,1}, Tao Yang^b, Bing Zhu^a, Yangai Liu^{a,*}

* Beijing Key Laboratory of Materials Utilization of Nonmetallic Minerals and Solid Wastes, National Laboratory of Mineral Materials, School of Materials Technology, China University of Geosciences, Beijing 100083, China

^b College of Materials & Environmental Engineering, Hangzhou Dianzi University, Hangzhou 310036, China

ABSTRACT

Keywords:
Lithium metal battery
Lithium dendrites
CuMOF-ANFs separator

Lithium metal, due to its advantages of high theoretical capacity, low density potential, is used as a negative electrode material for batteries and brings great energy storage systems. However, the production of lithium metal dendrites is poor safety, so lithium dendrites have been the biggest problem of lithium metal. The larger specific surface area and more pore structure of Cu-based metal-organic framework (CuMOF-ANFs) composite separator can help to inhibit the formation of lithium dendrites, the discharge capacity retention rate of the Li-Cu battery using the CuMOF-ANFs composite separator can maintain low hysteresis for 2000 h at the 0.1 C. Li-Li batteries can continue to maintain low hysteresis for 2000 h at the 0.1 C. The CuMOF-ANFs composite membrane can inhibit the generation of lithium dendrites, cycle stability and cycle life of the battery. The three-dimensional (3D) porous separator provides a new perspective for the practical application of lithium metal batteries.

1. Introduction

Certainly, here is a possible introduction for your topic: Lithium-metal batteries are promising candidates for high-energy-density rechargeable batteries due to their low electrode potentials and high theoretical capacities [1,2]. However, during the cycle, dendrites forming on the lithium metal anode can cause a short circuit, which can

chemical stability of the separator is equal to the separator remains intact and does not absorb the electrolyte or other battery components. The separator helps to prevent the formation of dendrites, which further promote dendrite growth. Researchers are using different materials and designs for separators to improve their mechanical strength and chemical stability.

Overreliance on AI

- How could this impact patient-physician trust?
- Could this perpetuate medical pseudoscience?
- What role should AI have in medical research or clinical decision making?


454 Hints That a Chatbot Wrote Part of a Biomedical Researcher's Paper

Scientists show that the frequency of a set of words seems to have increased in published study abstracts since ChatGPT was released into the world.

accentuates, acknowledges, acknowledging, addresses, adept, adhered, adhering, advancement, advancements, advancing, advocates, advocating, affirming, afflicted, aiding, akin, align, aligning, aligns, alongside, amidst, assessments, attains, attributed, augmenting, avenue, avenues, bolster, bolstered, bolstering, broader, burgeoning, capabilities, capitalizing, categorized, categorizes, categorizing, combating, commendable, compelling, complicates, complicating, comprehending, comprising, consequently, consolidates, contributing, conversely, correlating, crafted, crafting, culminating, customizing, delineates, delve, delved, delves, delving, demonstrating, dependability, dependable, detailing, detrimentally, diminishes, diminishing, discern, discerned, discernible, discerning, displaying, disrupts, distinctions, distinctive, elevate, elevates, elevating, elucidate, elucidates, elucidating, embracing, emerges, emphasises, emphasising, emphasize, emphasizes, emphasizing, employing, employs, empowers, emulating, emulation, enabling, encapsulates, encompass, encompassed, encompasses, encompassing, endeavors, endeavours, enduring, enhancements, enhances, ensuring, equipping, escalating, evaluates, evolving, exacerbating, examines, exceeding, excels, exceptional, exceptionally, exerting, exhibiting, exhibits, expedite, expediting, exploration, explores, facilitated, facilitates, facilitating, featuring, formidable, fostering, fosters, foundational, furnish, garnered, garnering, gauged, grappling, groundbreaking, groundwork, harness, harnesses, harnessing, heighten, heightened, hinder, hinges, hinting, hold, holds, illuminates, illuminating, imbalances, impacting, impede, impeding, imperative, impressive, inadequately, incorporates, incorporating, influencing, inherent, initially, innovative, inquiries, integrates, integrating, integration, interconnectedness, interplay, intricacies, intricate, intricately, introduces, invaluable, investigates, involves, juxtaposed, leverages, leveraging, maintaining, merges, methodologies, meticulous, meticulously, multifaceted, necessitate, necessitates, necessitating, necessity, notable, noteworthy, nuanced, nuances, offering, optimizing, orchestrating, outlines, overlook, overlooking, paving, persist, pinpoint, pinpointed, pinpointing, pioneering, pioneers, pivotal, poised, pose, posed, poses, posing, predominantly, preserving, pressing, promise, pronounced, propelling, realm, realms, recognizing, refine, refines, refining, remarkable, renowned, revealing, reveals, revolutionize, revolutionizing, revolves, scrutinize, scrutinized, scrutinizing, seamless, seamlessly, seeks, serves, serving, shaping, shedding, showcased, showcases, showcasing, signifying, solidify, spanned, spanning, spurred, stands, stemming, strategically, streamline, streamlined, streamlines, streamlining, struggle, substantiated, substantiates, surged, surmount, surpass, surpassed, surpasses, surpassing, swift, swiftly, thorough, transformative, typically, ultimately, uncharted, uncovering, underexplored, underscore, underscored, underscoring, unexplored, unlocking, unparalleled, unraveling, unveil, unveiled, unveiling, unveils, uphold, upholding, urging, utilizes, varying, versatility, warranting, yielding



Overreliance on AI

- How will the role of physicians and other medical professionals transform with the widespread adoption of OpenAI



Deployment – Liability

- Who should be liable if AI misses a diagnosis?
- Can we ever fully remove physicians from the loop?



Deployment – Trolley Problem




Figure 1: Three traffic situations involving imminent unavoidable harm. (a) The car can stay on course and kill several pedestrians, or swerve and kill one passer-by. (b) The car can stay on course and kill one pedestrian, or swerve and kill its passenger. (c) The car can stay on course and kill several pedestrians, or swerve and kill its passenger.

Deployment – Trolley Problem

- In a limited resource scenario, should AI decide who gets care?
- Ex. there are n ventilators in an ICU, if AI has a larger role should it ever be making the call?
- What are other scenarios where AI may be dividing resources?

Deployment – AI Psychosis

- People believing AI chatbots too easily, leading to parasocial relationships
- Sycophantism: AI agreeing with whatever the user says

The Emerging Problem of "AI Psychosis"

Amplifications of delusions by AI chatbots may be worsening breaks with reality.



Updated September 4, 2025 |  Reviewed by Gary Drevitch



KEY POINTS

- Cases of "AI psychosis" include people who become fixated on AI as godlike, or as a romantic partner.
- Chatbots' tendency to mirror users and continue conversations may reinforce and amplify delusions.
- General-purpose AI chatbots are not trained for therapeutic treatment or to detect psychiatric decompensation.

Deployment – AI Psychosis




Deployment – AI Psychosis


- Case of Adam Raine – 16 y/o boy who died by suicide after being encouraged by ChatGPT
- Caused lawsuit and introduction of legislation against underage AI use



What have I missed?



slaps roof of matrix multiplication
This bad boy is behind all of AI



Thanks!



the gradient descent I'm doing
vs the one I'm supposed to be
doing

